


Temporal attention block

*output sequence of BiLSTM-2
(64 forward + 64 backward units)*



$H = [h_1 \dots h_T]$, $T = 60$, h_t in $\mathbb{R}^{128}$

W : (128, 128) b : (128,) v : (128, 1) $\rightarrow$ 16,640 parameters

v is a learned query, so no external query is required and the block adds no recurrence.

Because α is a distribution over the T frames it is read directly as a per-frame saliency curve (Figure 7).